

Distributed Autonomous Systems M

Dante Piotto

spring semester 2024

Contents

1	Introduction and scenarios	5
1.1	Distributed Autonomous System	5
1.2	Scenarios and applications of distributed systems	5
1.3	Measurement filtering in wireless sensor networks	5
1.4	Parameter Estimation in Wireless Sensor Networks	6
1.4.1	Opinion Dynamics in Social Influence Networks	6
1.5	Main questions in averaging algorithms	6
1.6	Distributed control in cooperative robotics	6
1.6.1	Main questions in cooperative robotics	7
1.6.2	Distributed optimal control	7
2	Preliminaries on Algebraic Graph Theory	9
3	Averaging Systems	11
3.1	Distributed algorithm	11
3.2	Discrete-time averaging systems	11
3.3	Stochastic matrices	12
3.4	Example: Metropolis-Hastings weights	12
3.5	Time-varying digraphs	13
3.5.1	Averaging distributed algorithms over time-varying graphs	13
3.5.2	Discrete-time consensus over time-varying graphs	13
3.6	Laplacian dynamics	14
3.6.1	Properties of the Laplacian matrix	15
3.6.2	Consensus for Laplacian dynamics	15
4	Optimization basics	17
4.1	Optimization algorithms	18
4.1.1	A system theoretical perspective to the gradient method	18
4.2	steady-state analysis of the gradient method	18
4.2.1	Gradient method for quadratic programs	19
4.2.2	Gradient flow	19
4.2.3	Nesterov accelerated gradient method	19
5	Parallel Optimization and Federated Learning	21
5.1	Incremental gradient method	21
5.2	Stochastic Gradient Descent	22
5.2.1	Convergence	22
5.2.2	Stochastic mini-batch gradient method	22
5.3	Beyond SGD: Adaptive Momentum Estimation (Adam)	23
5.4	Federated learning	23
5.5	Distributed learning	23
6	Leader-Follower networks: Containment	25
6.1	Definition	25
6.2	Control law	25

7	Leader Follower networks: Formation control	27
7.1	analogy with mass-spring systems	27
8	Distributed Aggregative Optimization	31
8.1	Aggregative optimization	31
8.2	Distributed aggregative optimization	32
8.2.1	Extension to online aggregative optimization	33
9	Learning with Neural Networks	35
9.1	Basic element: the neuron model	35
10	Multi-Robot Safety Controllers	37
10.1	Control Barrier Functions	37
10.1.1	Admissible Control Space	37
10.1.2	Safety Filters via Control Barrier Certificates	37
10.2	Single-Robot Obstacle Avoidance: Single Integrator Model	38
10.3	Multi-Robot Collision Avoidance: Single Integrators	38
10.3.1	Control Barrier Function for Multi-Robot Collision Avoidance	38
10.3.2	Centralized Safety Controller for Multi-Robot Collision Avoidance	39
10.3.3	Decentralized Safety Controller for Multi-Robot Collision Avoidance	39
11	Stability Analysis for Linear Averaging	41
11.1	Irreducible and primitive matrices	41
11.2	Spectral properties of square matrices	41
12	Constraint-coupled Distributed Optimization	43
12.1	Constraint-coupled Setup	43
12.2	Multi-robot scenario	43
12.3	Distributed optimal control	44
12.4	Duality for constraint-coupled optimization	44
12.5	Centralized dual subgradient	45
12.5.1	Distributed dual subgradient convergence	46

Chapter 1

Introduction and scenarios

1.1 Distributed Autonomous System

Each agent $i \in \{1, \dots, N\}$ has

- local physical and/or cyber state x_i
- computational and sensing capabilities
- communication capability: exchange messages with "neighbours"

1.2 Scenarios and applications of distributed systems

- Averaging: distributed estimation, opinion dynamics
- Distributed control in cooperative robotics
- Distributed optimization
 - distributed machine learning
 - distributed decision-making in cooperative robotics
 - distributed optimal control in energy systems and cooperative robotics

1.3 Measurement filtering in wireless sensor networks

Consider a network of N sensors with local sensing, computation and communication. Agent $i, i \in \{1, \dots, N\}$, takes a local measurement from the environment (temperature, pressure, etc.). Let $x_{i0} \in \mathbb{R}$ be the scalar local measurement. Agents are interested in agreeing on the average of the measurements,

$$x_{\text{avg}} = \frac{1}{N} \sum_{i=1}^N x_{i0}$$

to have a better estimate of the environment quantity

Consider the following "distributed algorithm" based on "local" linear averaging, for each $i \in \{1, \dots, N\}$

$$\begin{aligned} x_i^0 &= x_{i0} \\ x_i^{k+1} &= \text{average}(x_i^k, \{x_j^k, j \text{ "neighbour" of } i\}), \quad k \in \mathbb{N} \end{aligned}$$

generalizing coefficients of the update:

$$\begin{aligned} x_i^0 &= x_{i0} \\ x_i^{k+1} &= \sum_{j=1}^N a_{ij} x_j^k \quad k \in \mathbb{N} \end{aligned}$$

Remark. $a_{ij} \geq 0$ and $\sum_{j=1}^N a_{ij} = 1$

Remark. $a_{ij} = 0$, for some $j \in \{1, \dots, N\}$, i.e. $a_{ij} = 0$ if i does not have access to the estimate of j

1.4 Parameter Estimation in Wireless Sensor Networks

Consider a network of N sensors with local sensing, computation and communication aiming at estimating a common parameter $\theta^* \in \mathbb{R}$. Each sensor i measures

$$y_i = B_i \theta^* + v_i$$

with $y_i \in \mathbb{R}^{m_i}$, B_i known matrix and v_i a random measurement noise. Assume v_1, \dots, v_N independent and Gaussian, with zero mean and covariance $E[v_i v_i^T] = \Sigma_i$. Assume $\sum_{i=1}^N m_i \geq m$ and $\begin{bmatrix} B_1 \\ \vdots \\ B_N \end{bmatrix}$ full rank. Compute a least-squares estimate

$$\hat{\theta} = \arg \min_{\theta} \sum_{i=1}^N (y_i - B_i \theta)^T \Sigma_i^{-1} (y_i - B_i \theta)$$

The optimal solution is

$$\begin{aligned} \hat{\theta} &= \left(\sum_{i=1}^N B_i^T \Sigma_i^{-1} B_i \right)^{-1} \sum_{i=1}^N B_i^T \Sigma_i^{-1} y_i \\ &= \left(\frac{1}{N} \sum_{i=1}^N B_i^T \Sigma_i^{-1} B_i \right)^{-1} \frac{1}{N} \sum_{i=1}^N B_i^T \Sigma_i^{-1} y_i \end{aligned}$$

The optimal solution can be obtained by computing two averages $\frac{1}{N} \sum_{i=1}^N \beta_i$ and $\frac{1}{N} \sum_{i=1}^N \beta_i$

1.4.1 Opinion Dynamics in Social Influence Networks

Group of N individuals, with x_i^k being the opinion of individual i at time k . Opinions are updated according to

$$x_i^{k+1} = \sum_{j=1}^N a_{ij} x_j^k$$

1.5 Main questions in averaging algorithms

- Do node estimates converge? Do they converge to a common value ("reach consensus")?
- Do they reach consensus to the average ("average consensus")?
- How can we model communication in general networks?
- Can we answer the above questions for general networks and communication protocols?
- What assumptions do we need on the communication network?

1.6 Distributed control in cooperative robotics

Team of N (mobile) robots aiming to execute complex tasks

Basic tasks

- rendezvous, containment
- formation, flocking
- coverage

Complex tasks

- pickup and delivery
- surveillance and patrolling
- exploration
- satellite constellation

1.6.1 Main questions in cooperative robotics

- Do robot states asymptotically converge?
- Do the asymptotic states satisfy the global, desired task?
- How can we model communication in (general) robotic networks?
- What assumptions do we need on the communication network?
- Can we answer the above questions for general networks and communication protocols?

1.6.2 Distributed optimal control

$$\begin{aligned}
& \min_{\substack{x_1, \dots, x_N \\ u_1, \dots, u_N}} \sum_{i=1}^N \left(\sum_{\tau=0}^{T-1} \ell_i(z_{i,\tau}, u_{i,\tau}) + m_i(z_{i,T}) \right) \\
& \text{subj to } \sum_{i=1}^N H_i z_{i,\tau} \leq h, & \tau \in [0, T] \\
& z_{i,\tau+1} = A_i z_{i,\tau} + B_i u_{i,\tau} & \forall i, \tau \in [0, T] \\
& z_{i,\tau} \in Z_i, \quad u_{i,\tau} \in U_i, & \forall i, \tau \in [0, T]
\end{aligned}$$

Chapter 2

Preliminaries on Algebraic Graph Theory

Definition 2.1 (Digraph)

A digraph is a pair $G = (I, E)$ where $I = 1, \dots, N$ is a set of elements called *nodes* and $E \subset I \times I$ is a set of ordered node pairs called *edges*

Edge: the pair (i, j) denotes an edge from i to j

Self-loop: edge from a node to itself, i.e. (i, i)

Definition 2.2 (Undirected (di)graph)

if for any $(i, j) \in E$ then $(j, i) \in E$

Definition 2.3 (Subgraph)

(I', E') subgraph of (I, E) if $I' \subset I$ and $E' \subset E$. Spanning subgraph if $I' = I$

Definition 2.4 (In-neighbours of i)

$j \in I$ is an in-neighbour of $i \in I$ if $(j, i) \in E$

Definition 2.5 (Set of in-neighbours of i)

$\mathcal{N}_i^{\text{IN}} = \{j \in \{1, \dots, N\} | (j, i) \in E\}$

Definition 2.6 (Out-neighbours of i)

$j \in I$ is an out-neighbour of $i \in I$ if $(i, j) \in E$

Definition 2.7 (Set of out-neighbours of i)

$\mathcal{N}_i^{\text{OUT}} = \{j \in \{1, \dots, N\} | (i, j) \in E\}$

Definition 2.8 (In-degree \deg_i^{IN})

number of in-neighbours, i.e. cardinality of $\mathcal{N}_i^{\text{IN}}$ ($\deg_i^{\text{IN}} = |\mathcal{N}_i^{\text{IN}}|$)

Out-degree analogous

Definition 2.9 (Balanced digraph)

A digraph G is balanced if $\deg_i^{\text{IN}} = \deg_i^{\text{OUT}}$ for all $i \in \{1, \dots, N\}$

Definition 2.10 (Directed path)

ordered sequence of nodes s.t. any pair of consecutive nodes is a directed edge. A path is *simple* if no node appears more than once

Definition 2.11 (Directed cycle)

simple directed path that starts and ends at the same node. Cycles of length one are called self-loops. Acyclic digraph if there are no cycles.

Definition 2.12 (Directed tree)

Acyclic digraph s.t. there exists a node, root, s.t. any node can be reached by only one directed path starting at the root.

Definition 2.13 (Spanning tree)

a spanning subgraph that is a (directed) tree

Definition 2.14 (Periodic graph)

if there exists a $k > 1$ period, that divides the length of every cycle.

Note: a graph with self-loops is aperiodic.

Definition 2.15 (Strongly connected digraph)

if there exists a directed path from any node to any other node

Definition 2.16 (Connected undirected graph)

if there exists a path from any node to any other node

Definition 2.17 (Globally reachable node)

if one of its nodes can be reached from any other node by traversing a directed path

Definition 2.18 (Complete graph)

Unweighted graph such that $\forall i, j \exists (i, j), (j, i) \in E$

Chapter 3

Averaging Systems

3.1 Distributed algorithm

Given a network of N agents communicating according to a fixed digraph G , i.e. each agent i can receive messages only from in-neighbours in the graph, i.e. from $j \in \mathcal{N}_i^{\text{IN}}$. We start by considering a fixed graph, thus, each agent communicates with the same neighbours at each iteration $k \in \mathbb{N}$

$$x_i^{k+1} = \text{stf}_i(x_i^k, \{x_j^k\}_{j \in \mathcal{N}_i^{\text{IN}}}), \quad i \in \{1, \dots, N\}$$

where stf_i is a function depending only on state x_i and states $x_j, j \in \mathcal{N}_i^{\text{IN}}$.

Alternative version with out-neighbours:

$$x_i^{k+1} = \text{stf}_i(\{x_j\}_{j \in \mathcal{N}_i^{\text{OUT}}})$$

3.2 Discrete-time averaging systems

Let $G^{\text{comm}} = (I, E)$ be a fixed (communication) digraph (self loops included). A linear averaging distributed algorithm can be written as:

$$x_i^{k+1} = \sum_{j \in \mathcal{N}_i^{\text{IN}}} a_{ij} x_j^k \quad i \in \{1, \dots, N\}$$

where $x_i^k \in \mathbb{R}$ is the state of agent i at k and $a_{ij} > 0$ are positive weights.

Remark. The weights a_{ij} are defined only for $(i, j) \in E$

Each i uses only the states of neighbours $j \in \mathcal{N}_i^{\text{IN}}$, thus distributed algorithm.

For analysis purposes, let us define weights $a_{ij} = 0$ for $(j, i) \notin E$. Thus we can rewrite the distributed algorithm as

$$x_i^{k+1} = \sum_{j=1}^N a_{ij} x_j^k \quad i \in \{1, \dots, N\}$$

This is a LTI autonomous system

$$\begin{bmatrix} x_1^{k+1} \\ \vdots \\ x_N^{k+1} \end{bmatrix} = \begin{bmatrix} a_{11} & \cdots & a_{1N} \\ \vdots & \ddots & \vdots \\ a_{N1} & \cdots & a_{NN} \end{bmatrix} \begin{bmatrix} x_1^k \\ \vdots \\ x_N^k \end{bmatrix}$$

Which can be compactly written as

$$x^{k+1} = A x^k$$

Remark. The matrix A can be seen as the weighted adjacency matrix of the reverse digraph $G^{\text{comm, rev}}$ of the digraph G^{comm}

If instead of in-neighbours we use out-neighbours, we call the digraph a sensing digraph G^{sens} . In this case the notation becomes consistent with graph theory, so we get

$$x^{k+1} = Ax^k$$

where A can be seen as the weighted adjacency matrix of the sensing digraph G^{sens}

3.3 Stochastic matrices

The non-negative square matrix $A \in \mathbb{R}^{N \times N}$ is

- row stochastic if $A\mathbf{1} = \mathbf{1}$ (each row sums to 1)
- column stochastic if $A^\top \mathbf{1} = \mathbf{1}$ (each column sums to 1)
- doubly stochastic if both row and column stochastic.

Lemma. Let A be a row-stochastic matrix and G the associate digraph. If G is strongly connected and aperiodic, then

1. the eigenvalue $\lambda = 1$ is simple;
2. all the other eigenvalues μ satisfy $|\mu| < 1$

Remark. The condition "G contains a globally reachable node and the subgraph of globally reachable nodes is aperiodic" is necessary and sufficient

Theorem 3.1 (Consensus)

Consider a (discrete-time) averaging system with associated digraph G and weighted adjacency matrix A . Assume G is strongly connected and aperiodic, and A is row stochastic. Then

1. there exists a left eigenvector $w \in \mathbb{R}^N, w > 0$ (i.e. with positive components $w_i > 0$ for all $i = 1, \dots, N$) such that

$$\lim_{k \rightarrow \infty} x^k = \mathbf{1} \frac{w^\top x^0}{w^\top \mathbf{1}} = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} \frac{\sum_{i=1}^N w_i x_i^0}{\sum_{i=1}^N w_i}$$

i.e., consensus is reached to $\frac{\sum_{i=1}^N w_i x_i^0}{\sum_{i=1}^N w_i}$

2. if additionally A is doubly stochastic, then

$$\lim_{k \rightarrow \infty} x^k = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix} \frac{\sum_{i=1}^N x_i^0}{N}$$

i.e., average consensus is reached

3.4 Example: Metropolis-Hastings weights

Given an undirected unweighted graph G with edge set E and degrees d_1, \dots, d_n

$$a_{ij} = \begin{cases} \frac{1}{1 + \max\{d_i, d_j\}} & \text{if } (i, j) \in E \text{ and } i \neq j \\ 1 - \sum_{h \in \mathcal{N}_i \setminus \{j\}} a_{ih} & \text{if } i = j \\ 0 & \text{otherwise} \end{cases}$$

Result: the matrix A is symmetric and doubly-stochastic.

3.5 Time-varying digraphs

A time-varying digraph is a sequence of digraphs $\{G(k)\}_{k \geq 0}$.

Remark. The main definitions of in/out neighbours, in/out degree, adjacency matrix can be generalized by considering time-varying versions, i.e. $\mathcal{N}_i^{\text{IN}}(k)$, $\mathcal{N}_i^{\text{OUT}}(k)$, $\deg_i^{\text{IN}}(k)$, $\deg_i^{\text{OUT}}(k)$, $A(k)$ associated to each graph $G(k)$. Connectivity requires new definitions as assuming each $G(k)$ to be connected is too conservative.

Definition 3.1 (Jointly strongly connected digraph)

if $\bigcup_{\tau=k}^{+\infty} G(\tau)$ is strongly connected $\forall k \geq 0$

Definition 3.2 (Uniformly jointly strongly connected (or B -strongly connected) digraph)

if there exists $B \in \mathbb{N}$ such that $\bigcup_{\tau=k}^{k+B} G(\tau)$ is strongly connected $\forall k \geq 0$

Remark. The graph can be disconnected at some time k .

3.5.1 Averaging distributed algorithms over time-varying graphs

Let $\{G(k)\}_{k \geq 0}$ be a time-varying digraph (with self loops for each $G(k)$). Consider the distributed algorithm

$$x_i^{k+1} = \sum_{j \in \mathcal{N}_i^{\text{IN}}(k)} a_{ij}(k) x_j^k \quad \forall i \in \{1, \dots, N\}$$

or the out-neighbours version

$$x_i^{k+1} = \sum_{j \in \mathcal{N}_i^{\text{OUT}}(k)} a_{ij}(k) x_j^k \quad \forall i \in \{1, \dots, N\}$$

where $x_i^k \in \mathbb{R}$ is the state of agent i at k and $a_{ij}(k) > 0$.

For analysis purposes, let us define weights $a_{ij}(k) = 0$ for $(i, j) \notin E(k)$. Thus we can rewrite the distributed algorithm as

$$x_i^{k+1} = \sum_{j=1}^N a_{ij}(k) x_j^k \quad i \in \{1, \dots, N\}$$

This is a Linear Time-Varying system

$$x^{k+1} = A(k)x^k$$

with state $x := [x_1, \dots, x_N]^\top$ and state matrix

$$A(k) := \begin{bmatrix} a_{11}(k) & \dots & a_{1N}(k) \\ \vdots & \ddots & \vdots \\ a_{N1}(k) & \dots & a_{NN}(k) \end{bmatrix}$$

being a weighted adjacency matrix associated to the digraph $G(k)$.

3.5.2 Discrete-time consensus over time-varying graphs

Theorem 3.2

Let $\{A(k)\}_{k \geq 0}$ be a sequence of row-stochastic matrices with associated digraphs $\{G(k)\}_{k \geq 0}$. Assume

1. each digraph $G(k)$ has a self-loop at each node;
2. each non-zero edge weight $a_{ij}(k)$, including the self-loop weights $a_{ii}(k)$, is larger than a constant $\epsilon > 0$;
3. there exists $B \in \mathbb{N}$ such that, for all times $k \geq 0$, the union digraph $G(k) \cup \dots \cup G(k+B)$ is strongly connected.

Then

1. there exists a non-negative vector $w \in \mathbb{R}^N$ such that the solution to $x^{k+1} = A(k)x^k$ converges (exponentially) to $\mathbf{1} \frac{w^\top x^0}{w^\top \mathbf{1}}$, i.e.

$$\lim_{k \rightarrow \infty} x^k = \mathbf{1} \left(\frac{w^\top x^0}{w^\top \mathbf{1}} \right)$$

2. if additionally each matrix in the sequence is doubly-stochastic, then

$$\lim_{k \rightarrow \infty} x^k = \mathbf{1} \frac{1}{N} \sum_{i=1}^N x_i^0$$

i.e., average consensus is achieved

3.6 Laplacian dynamics

Consider a network of dynamical systems with dynamics

$$\dot{x}(t) = u_i(t) \quad i \in \{1, \dots, N\}$$

with states $x_i \in \mathbb{R}$ and inputs $u_i \in \mathbb{R}$, communicating (or interacting) according to a digraph G . Consider a (distributed) "proportional" feedback control

$$u_i(t) = - \sum_{j \in \mathcal{N}_i^{\text{IN}}} a_{ij}(x_i(t) - x_j(t))$$

or the out-neighbour version

$$u_i(t) = - \sum_{j \in \mathcal{N}_i^{\text{OUT}}} a_{ij}(x_i(t) - x_j(t))$$

For analysis purposes, let us define weights $a_{ij}(k) = 0$ for $(i, j) \notin E(k)$. Thus we can rewrite the distributed control systems as

$$\dot{x}_i(t) = - \sum_{j=1}^N a_{ij}(x_i(t) - x_j(t)) \quad \forall i \in \{1, \dots, N\}$$

Defining $x := [x_1 \cdots x_N]^\top$, it can be shown that it can be rewritten as the following Linear Time Invariant continuous-time system

$$\dot{x}(t) = -Lx(t)$$

where L is the (weighted) Laplacian associated to the digraph G with (weighted) adjacency matrix A

Let

$$\dot{x}_i(t) = - \sum_{j=1}^N a_{ij}(x_i(t) - x_j(t)) \quad \forall i \in \{1, \dots, N\}$$

rearranging terms

$$\dot{x}_i(t) = - \left(\sum_{j=1}^N a_{ij} \right) x_i(t) + \sum_{j=1}^N a_{ij} x_j(t) = -\deg_i^{\text{OUT}} x_i(t) + (Ax(t))_i$$

where $(Ax(t))_i$ is the i -th element of $Ax(t)$. Writing the previous dynamics in a compact form

$$\dot{x}(t) = -(D^{\text{OUT}} - A)x(t)$$

where we recall that D^{OUT} is the (weighted) out-degree matrix. Recalling that $L = D^{\text{OUT}} - A$, it holds that

$$\dot{x}(t) = -Lx(t)$$

Remark. if the in-neighbours version is considered, then $\dot{x}(t) = -L^{\text{IN}}x(t)$, where $L^{\text{IN}} = D^{\text{IN}} - A^T$ is the in-degree Laplacian (i.e. the Laplacian of the reverse graph of G)

3.6.1 Properties of the Laplacian matrix

It can be easily verified that

$$L\mathbf{1} = D^{\text{OUT}}\mathbf{1} - A\mathbf{1} = \begin{bmatrix} \deg_1^{\text{OUT}} \\ \vdots \\ \deg_i^{\text{OUT}} \end{bmatrix} - \begin{bmatrix} \deg_1^{\text{OUT}} \\ \vdots \\ \deg_i^{\text{OUT}} \end{bmatrix} = 0$$

i.e., $\lambda = 0$ is an eigenvalue of L and $\mathbf{1}$ is an associated eigenvector.

Lemma. Given a weighted digraph with Laplacian L , then all eigenvalues of L different from zero have strictly positive real part

Lemma. Given a weighted digraph with Laplacian L , the following statements are equivalent:

1. G is weight-balanced, i.e. $D^{\text{IN}} = D^{\text{OUT}}$
2. $\mathbf{1}L = 0$

Theorem 3.3

A weighted digraph with Laplacian L contains a globally reachable node if and only if $\lambda = 0$ is simple.

Corollary. If a weighted digraph is strongly connected, then $\lambda = 0$ is simple

3.6.2 Consensus for Laplacian dynamics

Theorem 3.4

let L be a (weighted) Laplacian matrix with associated strongly connected (weighted) digraph G . Consider the Laplacian dynamics $\dot{x}(t) = -Lx(t)$, $t \geq 0$, then

- 1.

$$\lim_{t \rightarrow \infty} x(t) = \mathbf{1} \left(\frac{w^\top x(0)}{w^\top \mathbf{1}} \right)$$

with $w^\top L = 0$, i.e. w is a left eigenvector for the eigenvalue $\lambda = 0$;

2. if additionally G is weight-balanced then

$$\lim_{t \rightarrow \infty} x(t) = \mathbf{1} \frac{\sum_{i=1}^N x_i(0)}{N}$$

Chapter 4

Optimization basics

Convexity and gradient monotonicity

If a convex function ℓ is also differentiable, then its gradient $\nabla\ell : \mathbb{R}^d \rightarrow \mathbb{R}^d$ satisfies

$$(\nabla\ell(z_A) - \nabla\ell(z_B))^\top (z_A - z_B) \geq 0$$

for all z_A, z_B . That is, the gradient $\nabla\ell$ is a monotone operator

Strict convexity and gradient monotonicity

A function ℓ is strictly convex if for $z_A \neq z_B$ and $\theta \in (0, 1)$

$$\ell(\theta z_A + (1 - \theta)z_B) < \theta\ell(z_A) + (1 - \theta)\ell(z_B)$$

If the strictly convex function ℓ is also differentiable, then its gradient satisfies

$$(\nabla\ell(z_A) - \nabla\ell(z_B))^\top (z_A - z_B) > 0$$

for all z_A, z_B . That is, the gradient $\nabla\ell$ is a strictly monotone operator

Strong convexity and gradient monotonicity

A function ℓ is strongly convex with parameter $\mu > 0$ if for $z_A \neq z_B$ and $\theta \in (0, 1)$

$$\ell(\theta z_A + (1 - \theta)z_B) < \theta\ell(z_A) + (1 - \theta)\ell(z_B) - \mu\theta(1 - \theta)\|z_A - z_B\|^2$$

The gradient of a differentiable strongly convex function satisfies

$$(\nabla\ell(z_A) - \nabla\ell(z_B))^\top (z_A - z_B) \geq \mu\|z_A - z_B\|^2$$

for all z_A, z_B . That is, the gradient $\nabla\ell$ is a strongly monotone operator

Convexity and Lipschitz continuity of the gradient

Consider a differentiable convex function ℓ with a Lipschitz continuous gradient with parameter $L > 0$, i.e.

$$\|\nabla\ell(z_A) - \nabla\ell(z_B)\| \leq L\|z_A - z_B\|$$

for all z_A, z_B . Then, the following characterization holds

$$(\nabla\ell(z_A) - \nabla\ell(z_B))^\top (z_A - z_B) \geq \frac{1}{L}\|\nabla\ell(z_A) - \nabla\ell(z_B)\|^2$$

for all z_A, z_B . That is, the gradient $\nabla\ell$ is a co-coercive operator

Strong convexity and Lipschitz continuity of the gradient

Consider a strongly convex (with parameter $\mu > 0$) function ℓ with Lipschitz continuous gradient (with parameter $L > 0$). The the followin characterization holds

$$(\nabla\ell(z_A) - \nabla\ell(z_B))^\top (z_A - z_B) \geq \frac{\mu L}{\mu + L} \|z_A - z_B\|^2 + \frac{1}{\mu + L} \|\nabla\ell(z_A) - \nabla\ell(z_B)\|^2$$

for all z_A, z_B .

4.1 Optimization algorithms

We consider optimization algorithms based on iterative descent.

Notation. We denote by $z^k \in \mathbb{R}^d$ the estimate at iteration $k \in \mathbb{N}$ of a local minimum.

The algorithm starts at a given initial guess ecc ecc we know iterative descent

4.1.1 A system theoretical perspective to the gradient method

Let ℓ be μ -strongly convex and with L -Lipschitz continuous gradient. The gradient method is a discrete-time integrator in feedback interconnection with a static map

$$\begin{aligned} z^{k+1} &= z^k - \alpha u^k, & z^0 \text{ given} \\ u^k &= \nabla\ell(z^k) \end{aligned}$$

where u^k is a "control input" obtained through a static (nonlinear) state feedback. The block diagram is

This is known as the *Lur'e problem*

4.2 steady-state analysis of the gradient method

Assume that the state z^k converges to some value z_{eq} . Then, such an equilibrium must satisfy:

$$\begin{aligned} z_{eq} = z_{eq} - \alpha \nabla\ell(z_{eq}) &\implies & z_{eq} : \nabla\ell(z_{eq}) \\ &\implies & z_{eq} = z^* \end{aligned}$$

Consider the change of coordinates $z^k \rightarrow \tilde{z}^k - z_{eq} = z^k - z^*$. Then, the error dynamics is

$$\begin{aligned} \tilde{z}^{k+1} &= \tilde{z}^k - \alpha u^k \\ u^k &= \nabla\ell(\tilde{z}^k + z^*) - \nabla\ell(z^*) \end{aligned}$$

where u^k and \tilde{z}^k satisfy, in light of the assumption on ℓ , the following inequality¹

$$-(u^k)^\top \tilde{z}^k \leq -\gamma_1 \|\tilde{z}^k\|^2 - \gamma_2 \|u^k\|^2$$

Consider a Lyapunov function $V : \mathbb{R}^d \rightarrow \mathbb{R}_{\geq 0}$ given by $V(\tilde{z}) = \|\tilde{z}\|^2$. Then

$$\begin{aligned} V(\tilde{z}^{k+1}) - V(\tilde{z}^k) &= \|\tilde{z}^{k+1}\|^2 - \|\tilde{z}^k\|^2 \\ &= \|\tilde{z}^k\|^2 - 2\alpha(u^k)^\top \tilde{z}^k + \alpha^2 \|u^k\|^2 - \|\tilde{z}^k\|^2 \\ &\leq -2\alpha\gamma_1 \|\tilde{z}^k\|^2 + \alpha(\alpha - 2\gamma_2) \|u^k\|^2 \end{aligned}$$

For a small enough stepsize α (i.e., $\alpha \leq 2\gamma_2$), we can write

$$\begin{aligned} V(\tilde{z}^{k+1}) - V(\tilde{z}^k) &< -2\alpha\gamma_1 \|\tilde{z}^k\|^2 \implies \|\tilde{z}^{k+1}\|^2 \leq (1 - 2\alpha\gamma_1) \|\tilde{z}^k\|^2 \\ &\leq (1 - 2\alpha\gamma_1)^k \|\tilde{z}^0\|^2 \end{aligned}$$

Therefore $\{\tilde{z}^k\}_{k \in \mathbb{N}}$ goes exponentially/geometrically fast to zero

¹For all z_A, z_B it holds that

$$(\nabla\ell(z_A) - \nabla\ell(z_B))^\top (z_A - z_B) \geq \frac{\mu L}{\mu + L} \|z_A - z_B\|^2 + \frac{1}{\mu + L} \|\nabla\ell(z_A) - \nabla\ell(z_B)\|^2$$

4.2.1 Gradient method for quadratic programs

Consider a quadratic program

$$\min_z \frac{1}{2} z^\top Q z + r^\top z$$

With $Q = Q^\top > 0$ The gradient method is an affine linear system

$$\begin{aligned} z^{k+1} &= z^k - \alpha(Qz^k + r) \quad z^k \text{ given} \\ &= (I - \alpha Q)z^k - \alpha r \end{aligned}$$

For a sufficiently small α , the state matrix $(I - \alpha Q)$ is Schur. Hence, the state trajectory is ²

$$z^k = (I - \alpha Q)^k z^0 - \alpha \sum_{i=0}^{k-1} (I - \alpha Q)^i r \xrightarrow{k \rightarrow \infty} -\alpha \left(\sum_{i=0}^{\infty} (I - \alpha Q)^i \right) r = -Q^{-1} r$$

4.2.2 Gradient flow

Let us swap the roles of the plant (the static nonlinearity) and the controller (the integrator) We obtain the so-called gradient flow (continuous-time dynamics)

$$\dot{z}(t) = -\nabla \ell(z(t)) \quad z(0) = z_0$$

Remark. A solution to the ODE exists if the vector field is Lipschitz continuous

4.2.3 Nesterov accelerated gradient method

Consider the following two-step algorithm: for all $k \in \mathbb{N}$

$$\zeta^{k+1} = \zeta^k + \alpha_1(\zeta^k - \zeta^{k+1}) - \alpha_2 \nabla \ell(\zeta^k + \alpha_1(\zeta^k - \zeta^{k+1})), \quad \zeta^0, \zeta^{-1} \text{ given}$$

for some $\alpha_1, \alpha_2 > 0$. It admits the state-space representation More general updates can also be considered: for all $k \in \mathbb{N}$

$$\zeta^{k+1} = \zeta^k + \alpha_1(\zeta^k - \zeta^{k+1}) - \alpha_2 \nabla \ell(\zeta^k + \alpha_3(\zeta^k - \zeta^{k-1})), \quad \zeta^0, \zeta^{-1} \text{ given}$$

for some $\alpha_1, \alpha_2, \alpha_3 > 0$

²The geometric series $\sum_{i=0}^{\infty} \rho^i$ is equal to $(1 - \rho)^{-1}$ for all $\rho < 1$

Chapter 5

Parallel Optimization and Federated Learning

Cost-coupled optimization for learning

In learning applications, we usually consider optimization problems in the form

$$\min_{z \in \mathbb{R}^d} \sum_{i=1}^N \ell_i(z)$$

where, for all $i = 1, \dots, N$, the cost function $\ell_i : \mathbb{R}^d \rightarrow \mathbb{R}$ is local and private

(Batch) gradient method for learning

Consider the optimization problem

$$\min_z \sum_{i=1}^N \ell_i(z)$$

The (batch) gradient method is: for each iteration $k \in \mathbb{N}$

$$z^{k+1} = z^k - \alpha \sum_{i=1}^N \nabla \ell_i(z^k)$$

Remark. computation can be expensive

5.1 Incremental gradient method

Consider the optimization problem

$$\min_z \sum_{i=1}^N \ell_i(z)$$

Idea: rather than using the whole batch gradient at each $k \in \mathbb{N}$, just select one single "sample" per iteration.

The *incremental gradient method* is: for each iteration $k \in \mathbb{N}$

$$z^{k+1} = z^k - \alpha \nabla \ell_{i^k}(z^k)$$

where $i^k \in \{1, \dots, N\}$.

Two rules for choosing index i^k at iteration k :

- Cyclic rule: choose $i^k = 1, 2, \dots, N, 1, 2, \dots, N, \dots$
- Randomized rule: draw $i^k \in \{1, \dots, N\}$ at random¹, e.g according to a discrete uniform distribution

¹Each i^k is a realization of a discrete random variable $\mathcal{X} \sim \mathcal{U}\{1, N\}$ in which every one of the N possible outcome values has an equal probability

5.2 Stochastic Gradient Descent

consider the stochastic optimization problem

$$\min_z \mathbb{E}_{\mathcal{W}}[\ell(z, \mathcal{W})]$$

where $\mathbb{E}_{\mathcal{W}}[\cdot]$ denotes the expected value with respect to the random variable \mathcal{W} (possibly having an unknown probability distribution $p_{\mathcal{W}}(w)$)

Remark. for all z , also $\ell(z, \mathcal{W})$ is a random variable, whose probability distribution depends on $p_{\mathcal{W}}$ and ℓ . Moreover, the gradient $\nabla \ell(z, \mathcal{W})$ at each z is a random quantity

Assumption: There exists an oracle that, given a realization \bar{w} of \mathcal{W} , returns the corresponding realization of the gradient $\nabla \ell(\bar{z}, \bar{w})$ at any query point \bar{z}

The stochastic gradient descent is: for each iteration $k \in \mathbb{N}$ draw a realization w^k of \mathcal{W} and update

$$z^{k+1} = z^k - \alpha \nabla \ell(z^k, w^k)$$

5.2.1 Convergence

Proposition 5.1 (convergence with constant step-size)

Assume:

- ℓ is a μ -strongly convex function with L -Lipschitz continuous gradient (uniformly in its second argument)
- $\nabla \ell(z, \mathcal{W})$ is an unbiased estimatee of $\nabla_z \mathbb{E}_{\mathcal{W}}[\ell(z, \mathcal{W})]$
- $\|\nabla \ell(z, \mathcal{W})\| \leq M$ almost surely² for some $M > 0$

Let $\{z^k\}_{k \in \mathbb{N}}$ be a realized sequence of the SGD with stepsize $\alpha \leq 1/2\mu$. Then

$$\|z^k - z^*\|^2 \leq (1 - 2\mu\alpha)^k \left(\|z^0 - z^*\|^2 - \frac{\alpha M^2}{2\mu} \right) + \frac{\alpha M^2}{2\mu}$$

Analysis idea: Use the sequence of realizations w^0, w^1, \dots, w^k of \mathcal{W} to approximate the original stochastic problem as

$$\mathbb{E}_{\mathcal{W}}[\ell(z, \mathcal{W})] \approx \frac{1}{K} \sum_{k=1}^L \ell(z, w^k)$$

5.2.2 Stochastic mini-batch gradient method

At each iteration $k \in \mathbb{N}$, one can use a set of realizations of \mathcal{W} to statistically approximate the (random) gradient of ℓ at a given point z^k

Remark. The set of realizations is denoted by \mathcal{I}^k and is usually called minibatch

For each $k \in \mathbb{N}$, choose $\mathcal{I}^k \subset \{1, \dots, N\}$ with $|\mathcal{I}^k| \ll N$. The stochastic minibatch gradient method is

$$z^{k+1} = z^k - \alpha \sum_{i \in \mathcal{I}^k} \nabla \ell(z^k, w^i)$$

Remark. The minibatch can be chosen randomly or deterministically

Remark. This approach can also be applied to a "deterministic" optimization problem

²A sequence of random variables $\{\mathcal{X}_k\}_{k \in \mathbb{N}}$ converges almost surely to the rv \mathcal{X} if $\mathbb{P}(\lim_{k \rightarrow \infty} \mathcal{X}_k = \mathcal{X}) = 1$

5.3 Beyond SGD: Adaptive Momentum Estimation (Adam)

The ADAM algorithm reads as follows

- Mean and Variance (first momentum and second momentum)

$$\begin{aligned} m^{k+1} &= \beta_1 m^k + (1 - \beta_1) \nabla \ell(z^k, w^k), & \text{for some } \beta \in (0, 1) \\ v^{k+1} &= \beta_2 v^k + (1 - \beta_2) [\nabla \ell(z^k, w^k)]^2, & \text{for some } \beta_2 \in (0, 1) \end{aligned}$$

where the square operation $[\cdot]^2$ is meant component-wise

- Construct the descent direction

$$\begin{aligned} \hat{m} &= \frac{1}{1 - \beta_1^{k+1}} m^{k+1} \\ \hat{v} &= \frac{1}{1 - \beta_2^{k+1}} v^{k+1} \\ d^k &= -\frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}, \quad \text{for some } \epsilon > 0 \end{aligned}$$

where the division in the last equation is meant element-wise

- Update of the solution estimate

$$z^{k+1} = z^k + \alpha d^k$$

5.4 Federated learning

Consider the optimization problem

$$\min_z \sum_{i=1}^N \ell(z; \mathcal{D}^i, p^i)$$

Paradigm:

- local private data $\mathcal{D}^i = ([\mathcal{D}^i]_1, [\mathcal{D}^i]_2, \dots, [\mathcal{D}^i]_d)$ and p^i
- learn common parameters $z^* \in \mathbb{R}^d$ (common neural network)
- communication with a parameter server

5.5 Distributed learning

Consider the optimization problem

$$\min_z \sum_{i=1}^N \ell(z; \mathcal{D}^i, p^i)$$

Paradigm:

- local private data $\mathcal{D}^i = ([\mathcal{D}^i]_1, [\mathcal{D}^i]_2, \dots, [\mathcal{D}^i]_d)$ and p^i
- learn common parameters $z^* \in \mathbb{R}^d$ (common neural network)
- communication with neighbours only

Chapter 6

Leader-Follower networks: Containment

Consider a network of N agents communicating/interacting according to a fixed, undirected graph $G = (\{1, \dots, N\}, E)$. Let $x_i(t) \in \mathbb{R}^d$ be the state of agent i and for simplicity suppose $d = 1$, i.e. $x_i(t) \in \mathbb{R}$.

6.1 Definition

Suppose that agents are divided into two groups:

- N_f followers
- $N - N_f$ leaders

Define $x = \begin{bmatrix} x_f \\ x_l \end{bmatrix}$ with $x_f \in \mathbb{R}^{N_f}$ the followers' state and $x_l \in \mathbb{R}^{N-N_f}$ the leaders' state. Consistently, partition the Laplacian matrix as $L = \begin{bmatrix} L_f & L_{fl} \\ L_{FL}^\top & L_l \end{bmatrix}$. If leaders and followers run the same Laplacian-based distributed control law, then the Laplacian dynamics can be written as

$$\begin{bmatrix} \dot{x}_f(t) \\ \dot{x}_l(t) \end{bmatrix} = - \begin{bmatrix} L_f & L_{fl} \\ L_{FL}^\top & L_l \end{bmatrix} \begin{bmatrix} x_f(t) \\ x_l(t) \end{bmatrix}$$

6.2 Control law

Now, suppose that leaders and followers have different roles:

- N_f followers running Laplacian dynamics
- $N - N_f$ leaders being stationary

The distributed control system is given by

$$\begin{aligned} \dot{x}_i(t) &= - \sum_{j \in \mathcal{N}_i} a_{ij}(x_i(t) - x_j(t)) & \forall i \in \{1, \dots, N_f\} \\ \dot{x}_i(t) &= 0 & \forall i \in \{N_f + 1, \dots, N\} \end{aligned}$$

and in compact form

$$\begin{bmatrix} \dot{x}_f(t) \\ \dot{x}_l(t) \end{bmatrix} = - \begin{bmatrix} L_f & L_{fl} \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x_f(t) \\ x_l(t) \end{bmatrix}$$

The dynamics of the followers is given by

$$\dot{x}_f(t) = -L_f x_f(t) - L_{fl} x_l$$

where $x_l(t) = x_l(0) = x_l$

Chapter 7

Leader Follower networks: Formation control

7.1 analogy with mass-spring systems

Consider a platoon of N masses such that each mass i is connected with mass $i - 1$ and $i + 1$ through a spring with elastic constants respectively $a_{i-1,i} = a_{i,i-1} > 0$ and $a_{i+1,i} = a_{i,i+1} > 0$. Let $x_i \in \mathbb{R}$ be the position of mass i

The elastic force at mass i , $F_{e,i}(x)$ is given by

$$F_{e,i}(x) = -a_{i,i-1}(x_i - x_{i-1}) - a_{i,i+1}(x_i - x_{i+1})$$

For each spring, we can write the associated elastic force as the negative gradient of the elastic (potential) energy, so that

$$F_{e,i}(x) = -\frac{\partial}{\partial x_i} \left(\frac{1}{2} a_{i,i-1} \|x_i - x_{i-1}\|^2 + \frac{1}{2} a_{i,i+1} \|x_i - x_{i+1}\|^2 \right)$$

Let us suppose that each spring can be written as the parallel of two springs with elastic constants respectively $\frac{1}{2} a_{i,i-1} > 0$ and $\frac{1}{2} a_{i,i+1} > 0$.

Let $x_i \in \mathbb{R}$ be the position of mass i . The elastic force at mass i , $F_{e,i}(x)$ is given by

$$F_{e,i} = -\left(\frac{1}{2} a_{i,i-1} + \frac{1}{2} a_{i,i-1} \right) (x_i - x_{i-1}) - \left(\frac{1}{2} a_{i,i+1} + \frac{1}{2} a_{i,i+1} \right) (x_i - x_{i+1})$$

As before, we can write the elastic force as the negative gradient of the elastic (potential) energy, i.e.

$$F_{e,i}(x) = -\frac{\partial}{\partial x_i} \left(\frac{1}{2} \frac{a_{i,i-1}}{2} \|x_i - x_{i-1}\|^2 + \frac{1}{2} \frac{a_{i,i-1}}{2} \|x_i - x_{i-1}\|^2 + \frac{1}{2} \frac{a_{i,i+1}}{2} \|x_i - x_{i+1}\|^2 + \frac{1}{2} \frac{a_{i,i+1}}{2} \|x_i - x_{i+1}\|^2 \right)$$

The total elastic (potential) energy of the mass-spring system can be written as

$$\begin{aligned} V(x) &= \sum_{i=1}^N \sum_{j \in \mathcal{N}_i} \frac{1}{2} \frac{a_{i,j}}{2} \|x_i - x_j\|^2 \\ &= \sum_{i=1}^N \sum_{j \in \mathcal{N}_i} V_{ij}(x_i, x_j) \end{aligned}$$

where we have defined $\mathcal{N}_i := \{i - 1, i + 1\}$ and $V_{ij}(x_i, x_j) := \frac{1}{2} \frac{a_{i,j}}{2} \|x_i - x_j\|^2$.

Thus, the elastic force at mass i can be seen as the negative gradient of the energy, i.e.

$$\begin{aligned} F_{e,i}(x) &= -\frac{\partial}{\partial x_i} \sum_{j \in \mathcal{N}_i} (V_{ij}(x_i, x_j) + V_{ji}(x_j, x_i)) \\ &= -\frac{\partial}{\partial x_i} V(x) \end{aligned}$$

This formulation can be extended to more general systems in which masses are interconnected according to a topology described by an undirected graph $G = (\{1, \dots, N\}, E)$

By adding a damping term on each mass, the system dynamics can be written as

$$\begin{aligned}\dot{x}_i &= v_i \\ m_i \dot{v}_i &= -v_i - \frac{\partial}{\partial x_i} V(x) \quad \forall i \in \{1, \dots, N\}\end{aligned}$$

where we have considered the damping coefficient equal to one.

If we assume that masses are small, we may write

$$v_i \approx -\frac{\partial}{\partial x_i} V(x)$$

so that the dynamics may be approximated by the following first order dynamics

$$\dot{x}_i = -\frac{\partial}{\partial x_i} V(x) \quad \forall i \in \{1, \dots, N\}$$

Consider a network of N agents communicating/interacting according to a fixed, undirected graph G . Let $x_i(t) \in \mathbb{R}^d$ be the state of agent i . Let agents run a Laplacian dynamics

$$\dot{x}_i = -\sum_{j \in \mathcal{N}_i} a_{ij}(x_i - x_j) \quad \forall i \in \{1, \dots, N\}$$

We can rewrite it as

$$\dot{x}_i(t) = -\sum_{j \in \mathcal{N}_i} \frac{\partial}{\partial x_i} (V_{ij}(x_i, x_j) + V_{ji}(x_j, x_i)) \quad \forall i \in \{1, \dots, N\}$$

with $V_{ij}(x) = \frac{1}{2} \frac{a_{i,j}}{2} \|x_i - x_j\|^2$.

By recalling the definition of the total energy

$$V(x) = \sum_{i=1}^N \sum_{j \in \mathcal{N}_i} V_{ij}(x_i, x_j)$$

the Laplacian dynamics

$$\dot{x} = -Lx$$

can be seen as a "gradient flow", i.e.

$$\dot{x} = -\nabla V(x)$$

Thus, the consensus configuration can be seen as a stationary point of V . This idea can be extended to general potential functions and applied to distributed control systems.

Consider a network of N autonomous agents communicating/interacting according to a fixed, undirected graph. Let $x_i(t) \in \mathbb{R}^d$ be the state of agent i . Consider a global potential function defined as

$$V(x) = \sum_{i=1}^N \sum_{j \in \mathcal{N}_i} V_{ij}(x_i, x_j)$$

such that (local) minima of the potential correspond to desired configurations of the team. The gradient flow dynamics $\dot{x} = -\nabla V(x)$ turns out to be distributed. That is,

$$\dot{x}_i(t) = \sum_{j \in \mathcal{N}_i} \frac{\partial}{\partial x_i} (V_{ij}(x_i, x_j) + V_{ji}(x_j, x_i)) \quad \forall i \in \{1, \dots, N\}$$

We can define a desired formation by assigning a set of distances, d_{ij} , between neighbouring agents i and j in a suitable graph

The main idea for formation control is to define a potential function matching the sparsity of G , $V^{\text{form}}(x) = \sum_{i=1}^N \sum_{j \in \mathcal{N}_i} V_{ij}^{\text{form}}(i, j)$, such that a configuration x^{teform} satisfying

$$\|x_i^{\text{form}} - x_j^{\text{form}}\| = d_{ij} \quad \forall (i, j) \in E$$

is a minimum of V .

In order to reach a formation with assigned distances d_{ij} , let us define

$$V_{ij}^{\text{form}}(x) = \frac{1}{8} (\|x_i - x_j\|^2 - d_{ij}^2)^2$$

with corresponding (global) potential function $V^{\text{form}}(x) = \sum_{i=1}^N \sum_{j \in \mathcal{N}_i} V_{ij}^{\text{form}}(x_i, x_j)$.

The gradient flow dynamics of each agent i is given by

$$\dot{x}_i = - \sum_{j \in \mathcal{N}_i} \frac{\partial}{\partial x_i} (V_{ij}^{\text{form}}(x_i, x_j) + V_{ji}(x_j, x_i)) \quad \forall i \in \{1, \dots, N\}$$

which reads as

$$\dot{x}_i = - \sum_{j \in \mathcal{N}_i} (\|x_i - x_j\|^2 - d_{ij}^2) (x_i - x_j) \quad \forall i \in \{1, \dots, N\}$$

This dynamics has multiple equilibrium points, including the desired formation in which the agents are at the assigned distances. In particular, the consensual solution $x_1 = x_2 = \dots = x_N$ is an (undesired) equilibrium.

Such a "degenerate" equilibrium can be avoided by means of additional "collision avoidance" potential functions $V_{ij}^{\text{ca}}(x_i, x_j)$ such that

$$\lim_{\|x_i - x_j\| \rightarrow 0} V_{ij}^{\text{ca}}(x_i, x_j) = +\infty$$

A possible solution is a barrier function given

$$V_{ij}^{\text{ca}} = -\log(\|x_i - x_j\|)$$

Similarly, barrier potential functions $V^{\text{obs}}(x_i)$, depending only on the state of agent x_i , can be used to avoid obstacles.

The formation control dynamics becomes

$$\dot{x}_i = - \frac{\partial V^{\text{form}}}{\partial x_i} - \frac{\partial V^{\text{ca}}(x)}{\partial x_i} - \nabla V^{\text{obs}}(x) \quad \forall i \in \{1, \dots, N\}$$

Chapter 8

Distributed Aggregative Optimization

Consider N robots in the plane that want to optimize their positions $z_i \in \mathbb{R}^2$, for all $i = 1, \dots, N$ to perform multi-robot surveillance. Let:

- $r_0 \in \mathbb{R}^2$ be a target to protect
- $r_i \in \mathbb{R}^2$ be the intruder associated to robot i
- $\sigma(z) = \frac{1}{N} \sum_{i=1}^N z_i$ is the barycenter of the robots
- Local cost function of robot i

$$\ell_i(z_i, \sigma(z)) = \gamma_i \|z_i - r_i\|^2 + \|\sigma(z) - r_0\|^2$$

with $z \in \mathbb{R}^{2N}$ the stack of z_1, \dots, z_N and $\gamma_i > 0$ being a tradeoff parameter.

8.1 Aggregative optimization

Let us consider aggregative optimization problems in the form

$$\min_{z_1, \dots, z_N} \sum_{i=1}^N \ell_i(z_i, \sigma(z))$$

where the aggregative variable $\sigma(z)$ is defined as

$$\sigma(z) = \frac{1}{N} \sum_{i=1}^N \phi_i(z_i)$$

where

- $z = (z_1, \dots, z_N)$, with each $z_i \in \mathbb{R}^{n_i}$, for all $i = 1, \dots, N$
- $\ell_i : \mathbb{R}^{n_i} \times \mathbb{R}^d \rightarrow \mathbb{R}$ and $\phi_i : \mathbb{R}^{n_i} \rightarrow \mathbb{R}^d$, for all $i = 1, \dots, N$

For scalar states, $z_i \in \mathbb{R}$, the *centralized gradient* method at iteration k reads as

$$z_i^{k+1} = z_i^k - \alpha \frac{\partial}{\partial z_i} \left(\sum_{j=1}^N \ell_j(z_j, \sigma(z_1, \dots, z_N)) \right) \Big|_{z_1=z_1^k, \dots, z_N=z_N^k}$$

for all $i = 1, \dots, N$ where $\alpha > 0$ is the stepsize

Gradient computation (scalar case)

Since the cost function is a composite function, we need the chain rule to compute its derivative with respect to z_i

$$\begin{aligned} & \left. \frac{\partial}{\partial z_i} \left(\sum_{j=1}^N \ell_j(z_j, \sigma(z_1, \dots, z_N)) \right) \right|_{z_1=z_1^k, \dots, z_N=z_N^k} \\ &= \left. \frac{\partial}{\partial z_i} \ell_i(z_i, \sigma) \right|_{z_i=z_i^k, \sigma=\frac{1}{N} \sum_{j=1}^N \phi_j(z_j^k)} \\ &+ \sum_{j=1}^N \left. \frac{\partial}{\partial \sigma} \ell_j(z_j, \sigma) \right|_{z_j=z_j^k, \sigma=\frac{1}{N} \sum_{j=1}^N \phi_j(z_j^k)} \cdot \left. \frac{\partial \sigma(z_1, \dots, z_N)}{\partial z_i} \right|_{z_1=z_1^k, \dots, z_N=z_N^k} \end{aligned}$$

Notice that $\frac{\partial \sigma(z_1, \dots, z_N)}{\partial z_i} = \frac{1}{N} \frac{d}{dz_i} \phi_i(z_i)$ can be computed locally

As in the scalar case, we use the chain rule to compute the gradient of the composite function. The i -th block of the gradient, denoted as $\left[\nabla \left(\sum_{j=1}^N \ell_j(z_j, \sigma(z_1, \dots, z_N)) \right) \right]_i \in \mathbb{R}^{n_i}$, is given by

$$\begin{aligned} & \left[\nabla \left(\sum_{j=1}^N \ell_j(z_j, \sigma(z_1, \dots, z_N)) \right) \right]_i \\ &= \nabla_1 \ell_i(z_i, \sigma) \Big|_{z_j=z_j^k, \sigma=\frac{1}{N} \sum_{j=1}^N \phi_j(z_j^k)} \\ &+ \frac{1}{N} \nabla \phi_i(z_i) \Big|_{z_i=z_i^k} \cdot \sum_{j=1}^N \nabla_2 \ell_j(z_j, \sigma) \Big|_{z_j=z_j^k, \sigma=\frac{1}{N} \sum_{j=1}^N \phi_j(z_j^k)} \end{aligned}$$

8.2 Distributed aggregative optimization

In a distributed context, each agent i

- knows only ℓ_i and ϕ_i
- maintains an estimate z_i^k of z_i^*
- maintains an estimate s_i^k of $\phi(z^k) = \frac{1}{N} \sum_{j=1}^N \phi_j(z_j^k)$
- maintains an estimate v_i^k of $\sum_{j=1}^N \nabla_2 \ell_j(z_j^k, \sigma(z^k))$

The "tracking" idea of gradient tracking algorithm is applied to aggregative optimization

$$\begin{aligned} z_i^{k+1} &= z_i^k - \alpha (\nabla_1 \ell_i(z_i^k, s_i^k) + \nabla \phi_i(z_i^k) v_i^k) & z_i^0 &\in \mathbb{R}^{n_i} \\ s_i^{k+1} &= \sum_{j \in \mathcal{N}_i} a_{ij} s_j^k + \phi_i(z_i^{k+1}) - \phi_i(z_i^k) & s_i^0 &= \phi_i(z_i^0) \\ v_i^{k+1} &= \sum_{j \in \mathcal{N}_i} a_{ij} v_j^k + \nabla_2 \ell_i(z_i^{k+1}, s_i^{k+1}) - \nabla_2 \ell_i(z_i^k, s_i^k) & v_i^0 &= \nabla_2 \ell_i(z_i^0, s_i^0) \end{aligned}$$

Theorem 8.1 (aggregative tracking distributed optimization algorithm: convergence)

Assume G is a strongly connected and aperiodic digraph, and A is doubly stochastic. Assume that each function ℓ_i is strongly convex, the gradients $\nabla_1 \ell_i$ and $\nabla_2 \ell_i$ are Lipschitz continuous, and ϕ_i is differentiable and Lipschitz continuous.

Then, there exists α^* such that for all $\alpha \in (0, \alpha^*)$ the sequences of local solution estimates $\{z_1^k, \dots, z_N^k\}_{k \in \mathbb{N}}$ generated by the aggregative tracking distributed optimization algorithm satisfy

$$\lim_{k \rightarrow \infty} \|z_i^k - z_i^*\| = 0$$

at a linear rate, for all $i = 1, \dots, N$

8.2.1 Extension to online aggregative optimization

Consider a time-varying instance of the problem

$$\begin{aligned} \min_z \quad & \sum_{i=1}^N \ell_i^k(z_i, \sigma^k(z)) \\ \text{subj. to } & z_i \in Z_i^k, \quad \forall i = 1, \dots, N \end{aligned}$$

where $Z_i^k \subset \mathbb{R}^{n_i}$ are local constraint sets.

The goal is to design an algorithm generating a sequence $\{z_i^k\}_{k \in \mathbb{N}}$ that "tracks" the solution $z^{k,*} = (z_1^{k,*}, \dots, z_N^{k,*})$ of the k -th problem instance.

Remark. A regret R_K can be introduced for the analysis. Under suitable assumptions, it can be proven that

Chapter 9

Learning with Neural Networks

In supervised learning we have labeled data whose generic element(sample) is composed by

- an input point $\mathcal{D} \in \mathbb{R}^d$, i.e. a vector $\mathcal{D} = ([\mathcal{D}]_1, \dots, [\mathcal{D}]_d)$
- a label $p \in \mathbb{R}$ associated to the input point

A dataset is usually made by \mathcal{M} samples, i.e. $((\mathcal{D}^1, p^1), \dots, (\mathcal{D}^{\mathcal{M}}, p^{\mathcal{M}}))$

Goal: reconstruct the unknown input-output map between inputs and labels

Strategy: approximate the map as a nonlinear function $\phi(\cong; \cdot) : \mathbb{R}^d \rightarrow \mathbb{R}$ parametrized by \cong

Problem: find the best parameters \cong^* based on data

9.1 Basic element: the neuron model

A generic neuron h is a computational unit that

- has a set of weights $u_h \in \mathbb{R}^d$
- elaborates a vector $x \in \mathbb{R}^d$
- computes a scalar quantity

$$x_h^+ = \phi(x^\top u_h)$$

with $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ being the *activation function*

Chapter 10

Multi-Robot Safety Controllers

Consider a nonlinear system

$$\dot{x}(t) = f(x(t), u(t)), \quad x \in \mathbb{R}^n \quad u \in U \subset \mathbb{R}^m$$

Let us define a safe (state) set

$$X^s := \{x \in \mathbb{R}^n | V^s(x) \geq 0\}$$

for some sufficiently regular function $V^s : \mathbb{R}^n \rightarrow \mathbb{R}$

Goal: design a feedback control law $\kappa(x), \kappa : \mathbb{R}^n \rightarrow \mathbb{R}^m$ such that the set X^s is forward invariant

10.1 Control Barrier Functions

Consider the time derivative of $V^s(x(t))$ along the system trajectories

$$\begin{aligned} \frac{d}{dt} V^s(x(t)) &= \nabla V^s(x(t))^\top f(x(t)) + \sum_{h=1}^m \nabla V^s(x(t))^\top g_h(x(t)) u_h(x(t)) \\ &= L_f V^s(x(t)) + L_g V^s(x(t)) u(t) \end{aligned}$$

where L_f and L_g represent Lie derivatives. We say that V^s is a Control Barrier Function (CBF) if there exists a continuous, strictly increasing function $\gamma : \mathbb{R} \rightarrow \mathbb{R}$, with $\gamma(0) = 0$, such that

$$\sup_{u \in U} \{L_f V^s(x) + L_g V^s(x) u + \gamma(V^s(x))\} \geq 0$$

for all $x \in \mathbb{R}^n$.

The above inequality induced by a CBF is called a *control barrier certificate*.

10.1.1 Admissible Control Space

We can define the set of admissible (safe) controllers for a given state x as

$$U^s(x) = \{u \in \mathbb{R}^m | L_f V^s(x) + L_g V^s(x) u + \gamma(V^s(x)) \geq 0\}$$

i.e., the set of inputs satisfying the control barrier certificates.

Remark. It can be shown that (sufficiently regular) feedback controllers satisfying the above condition render X^s forward invariant and asymptotically stable.

10.1.2 Safety Filters via Control Barrier Certificates

Let $u^{\text{ref}} \in \mathbb{R}^m$ be a (possibly unsafe) reference input, e.g., obtained by a higher-level controller. The safety controller is designed to be "minimally invasive", i.e. it aims at altering the reference controller as little as possible. A safe policy $u = \kappa(x)$ can be designed as

$$\begin{aligned} \kappa(x) &= \arg \min_{u \in \mathbb{R}^m} \|u - u^{\text{ref}}(x)\|^2 \\ \text{subj. to } &-L_f V^s(x) - L_g V^s(x) u - \gamma(V^s(x)) \leq 0 \end{aligned}$$

10.2 Single-Robot Obstacle Avoidance: Single Integrator Model

Consider a robot modeled as a single integrator, i.e.

$$\dot{x} = u$$

with $x \in \mathbb{R}^d$ being the state of the robot and $u \in \mathbb{R}^d$ the control input. The goal is to keep the robot at a safety distance d from an obstacle.

Consider a CBF defined as

$$V^s(x) = \|x - x_{extobs}\|^2 - d^2$$

where d denotes the safety distance to keep from the obstacle. The gradient wrt x is given by

$$\nabla V^s(x) = 2(x - x_{obs})$$

The safety, CBF-based policy can be computed as

$$\begin{aligned} \kappa_s(x) = \arg \min_{u \in \mathbb{R}^d} & \|u - u^{\text{ref}}(x)\|^2 \\ \text{subj. to} & -2(x - x_{obs})^\top u - \gamma(\|x - x_{obs} - d^2\|) \leq 0 \end{aligned}$$

Remark. For a given, fixed $x \in \mathbb{R}^d$, the optimization problem above is a quadratic program

Remark. The safety controller is meant to prevent collisions while being "minimally invasive" in the sense that it aims at altering the reference controller (coming from a higher level) as little as possible

Remark. From a practical viewpoint the above problem requires a computation time, while the safety policy should be available in continuous time

10.3 Multi-Robot Collision Avoidance: Single Integrators

Consider a team of N robots modeled as single integrators, i.e.

$$\dot{x}_i = u_i$$

with $x_i \in \mathbb{R}^d$ being the state of robot i and $u_i \in \mathbb{R}^d$ being the control input. The goal is to keep the robots at a safety distance greater than d .

To keep the notation simpler, we consider a scalar robot state, and input, i.e. $x_i, u_i \in \mathbb{R}$.

10.3.1 Control Barrier Function for Multi-Robot Collision Avoidance

Consider a "pair-wise", local CBF defined as

$$V_{i,j}^s(x_i, x_j) = \|x_i - x_j\|^2 - d^2$$

where d denotes the safety distance between robots i and j .

The gradient wrt x_i and x_j is given by

$$\begin{aligned} \nabla_1 V_{i,j}^s(x_i, x_j) &= 2(x_i - x_j) \\ \nabla_2 V_{i,j}^s(x_i, x_j) &= 2(x_j - x_i) \end{aligned}$$

while the gradient wrt all the other states is zero.

Let $u \in \mathbb{R}^N$ denote the stack of inputs of the N robots. The admissible control space for collision avoidance is obtained by considering the intersection of all the constraint sets related to a pair-wise control barrier certificates. Thus,

$$U^s(x) \{u \in \mathbb{R}^N \mid -\nabla_1 V_{i,j}^s(x_i, x_j)^\top u_i - \nabla_2 V_{i,j}^s(x_i, x_j) u_j - \gamma V_{i,j}^s(x_i, x_j) \leq 0, \forall j \in \mathcal{N}_i, \forall i \in \{1, \dots, N\}\}$$

where \mathcal{N}_i includes all the robots j that may collide with robot i .

Remark. If there are no apriori restrictions on the motion of the robots and they share the same area, we have $\mathcal{N}_i = \{1, \dots, N\}$. However, for some state configurations \mathcal{N}_i may involve only a restricted number of neighbouring robots.

10.3.2 Centralized Safety Controller for Multi-Robot Collision Avoidance

The safety, CBF-based policy is obtained by the following optimization problem

$$\begin{aligned} \kappa(x) = \arg \min_{u \in \mathbb{R}^N} & \sum_{i=1}^N \|u_i - u_i^{\text{ref}}\|^2 \\ \text{subj. to} & \quad -\nabla_1 V_{i,j}^s(x_i, x_j)^\top u_i - \nabla_2 V_{i,j}^s(x_i, x_j) u_j - \gamma V_{i,j}^s(x_i, x_j) \leq 0 \\ & \quad \|u_i\| \leq u_i^{\max} \\ & \quad \forall j \in \mathcal{N}_i, \forall i \in \{1, \dots, N\} \end{aligned}$$

where $u_1^{\text{ref}}, \dots, u_N^{\text{ref}}$ are reference inputs for each robot coming from a higher control layer, while $u_1^{\max}, \dots, u_N^{\max}$ are input bounds for each robot.

10.3.3 Decentralized Safety Controller for Multi-Robot Collision Avoidance

In cooperative (distributed) multi-robot scenarios it is desirable not to have a central coordinator. A more constrained version of the problem can be formulated to obtain a decentralized solver.

Safety controller for robot i :

$$\begin{aligned} \kappa(x) = \arg \min_{u \in \mathbb{R}^N} & \sum_{i=1}^N \|u_i - u_i^{\text{ref}}\|^2 \\ \text{subj. to} & \quad -2\nabla_1 V_{i,j}^s(x_i, x_j)^\top u_i - \frac{1}{2}\gamma V_{i,j}^s(x_i, x_j) \leq 0 \\ & \quad \|u_i\| \leq u_i^{\max} \\ & \quad \forall j \in \mathcal{N}_i \end{aligned}$$

Remark. The decentralized safety controller requires robot i to receive information by all other robots that may incur in a collision.

Remark. More complex models may be considered to take into account dynamics constraints

Remark. When more complex dynamics are considered together with input bounds, the (centralized or decentralized) optimization problem may become infeasible (safety cannot be guaranteed). More sophisticated approaches are required to synthesize safety barrier certificates with guaranteed feasibility.

Remark. When the objectives of multiple robots conflict with the safety barrier certificates, robots might incur into a deadlock, i.e., a configuration that is safe but does not allow for task completion.

Chapter 11

Stability Analysis for Linear Averaging

Consider N agents in a network modeled as a graph G with associated weighted adjacency matrix A
 Linear averaging distributed algorithm

$$x_i^{k+1} = \sum_{j \in \mathcal{N}_i} a_{ij} x_j^k, \quad x^0 \text{ is given} \quad i = 1, \dots, N$$

where $x_i^k \in \mathbb{R}$ is the state of agent i at time k and $a_{ij} \geq 0$

Each agent i uses only the states of neighbours $j \in \mathcal{N}_i$, thus it is a distributed algorithm. The goal is to reach consensus.

The linear averaging is

We want to study the stability properties of the equilibrium point(s) of such an autonomous system

The compact version of the linear averaging is an LTI dynamical system

$$x^{k+1} = Ax^k \quad x^0 \text{ given}$$

The equilibrium points can be computed as

$$x_{\text{eq}} = Ax_{\text{eq}} \implies (I - A)x_{\text{eq}} = 0$$

11.1 Irreducible and primitive matrices

A square matrix $A \in \mathbb{R}^{N \times N}$ is

- non-negative if all its entries are greater than or equal to zero and we write $A \geq 0$
- irreducible if $\sum_{k=0}^{N-1} A^k > 0$
 Equivalently, there exists no permutation matrix T such that TAT^\top is block triangular
- primitive if there exists an integer $k \in \{1, \dots, N\}$ such that A^k is positive
- positive¹ if all its entries are greater than zero and we write $A > 0$

11.2 Spectral properties of square matrices

The spectrum of a square matrix A , $\text{spec}(A)$, is the set of eigenvalues of A . The spectral radius, $\rho(A)$, is the maximum norm of the eigenvalues of A .

Theorem 11.1 (Gershgorin Theorem)

For any (real or complex) square matrix A , it holds

$$\text{spec}(A) \subset \bigcup_{i=1}^N \left\{ s \in \mathbb{C} \mid |s - a_{ii}| \leq \sum_{j=1, j \neq i}^N |a_{ij}| \right\}$$

¹Not to be confused with a *positive definite* matrix, denoted as $A \succ 0$

That is, the union of disks in \mathbb{C} centered at each diagonal entry a_{ii} with radius $\sum_{j=1, j \neq i}^N |a_{ij}|$

Theorem 11.2

Let G be a weighted digraph with $N \geq 2$ nodes and with weighted adjacency matrix A . Then A is irreducible iff G is strongly connected

Theorem 11.3

Let G be a weighted digraph with $N \geq 2$ nodes and with weighted adjacency matrix A . Then A is primitive iff G is strongly connected and aperiodic

Theorem 11.4 (Perron-Frobenius theorem)

Let $A \in \mathbb{R}^{N \times N}$ with $N \geq 2$ be a non-negative matrix, then

1. there exists a real eigenvalue $\lambda \geq |\mu| \geq 0$ for all other eigenvalues $\mu \in \mathbb{C}$ of A
2. the right eigenvector $v \in \mathbb{R}^N$ and left eigenvector $w \in \mathbb{R}^N$ associated to λ can be selected non-negative
If, additionally $A \in \mathbb{R}^{N \times N}$ is irreducible, then
3. the eigenvalue $\lambda \in \mathbb{R}$ is strictly positive and simple
4. the right eigenvector $v \in \mathbb{R}^N$ and left eigenvector $w \in \mathbb{R}^N$ associated to λ are unique and positive, possibly up to rescaling
If, additionally $A \in \mathbb{R}^{N \times N}$ is primitive, then
5. The eigenvalue $\lambda > |\mu|$ for all other eigenvalues $\mu \in \mathbb{C}$ of A

Chapter 12

Constraint-coupled Distributed Optimization

Constraint-coupled optimization problem

$$\begin{aligned} \min_{z_1, \dots, z_N} \quad & \sum_{i=1}^N \ell_i(z_i) \\ \text{subj. to} \quad & z_i \in Z_i \subset \mathbb{R}^{n_i} \quad i = 1, \dots, N \\ & \sum_{i=1}^N g_{ij}(z_i) \leq 0 \quad j = 1, \dots, r \end{aligned}$$

with $\ell_i : \mathbb{R}^{n_i} \rightarrow \mathbb{R}$ convex, $g_{ij} : \mathbb{R}^{n_i} \rightarrow \mathbb{R}$ convex and $Z_i \subset \mathbb{R}^{n_i}$ convex and compact.

Constraint-coupled optimization problem (vector form)

$$\begin{aligned} \min_{z_1, \dots, z_N} \quad & \sum_{i=1}^N \ell_i(z_i) \\ \text{subj. to} \quad & z_i \in Z_i \subset \mathbb{R}^{n_i} \quad i = 1, \dots, N \\ & \sum_{i=1}^N g_i(z_i) \leq 0 \end{aligned}$$

with $\ell_i : \mathbb{R}^{n_i} \rightarrow \mathbb{R}$ convex, $g_i : \mathbb{R}^{n_i} \rightarrow \mathbb{R}^r$ convex and $Z_i \subset \mathbb{R}^{n_i}$ convex and compact

12.1 Constraint-coupled Setup

- N agents communicate over graph G
- (z_1, \dots, z_N) decision variables - size grows with N
- agent i knows ℓ_i, g_i and Z_i

12.2 Multi-robot scenario

Let us assume a framework with N robots that must collectively serve N tasks. We model the state of the i -th robot as $z_i \in \mathbb{R}^{n_i}$ where n_i is the number of tasks robot i can serve, and we denote

$$\begin{cases} z_{ik} = 1 & \text{if robot } i \text{ serves task } k \\ z_{ik} = 0 & \text{otherwise} \end{cases}$$

And we say that robot i serves task j with some cost c_{ij} . The constraint set is thus

$$Z_i = \{z_i \in \mathbb{R}^{n_i} | z_i \in \{0, 1\}, \sum_{k=1}^N z_{ik} = 1\}$$

Suppose each robot can serve all tasks. then

$$\sum_{i=1}^N z_{ik} = 1 \quad k = 1, \dots, N$$

in a more compact way

$$\sum_{i=1}^N z_i = 1$$

if robots can serve only a subset of tasks we can define some "selection" matrices $H_i, i \in \{1, \dots, N\}$. We can thus write the constraint as

$$\sum_{i=1}^N H_i z_i = 1$$

To summarize, we can rewrite the problem as

$$\begin{aligned} \min_{z_1, \dots, z_N} \quad & \sum_{i=1}^N c_i^T z_i \\ \text{subj. to } \quad & z_i \in Z_i \quad i = 1, \dots, N \\ & \sum_{i=1}^N H_i z_i = 1 \end{aligned}$$

Remark. Task allocation problems satisfy the "unimodularity" property. We can relax to a linear program (in place of an integer program):

$$Z_i = \{z_i \in \mathbb{R}^{n_i} | 0 \leq z_i \leq 1, \sum_{k=1}^N z_{ik} = 1\}$$

The solution $z_i^* \in \{0, 1\}$

12.3 Distributed optimal control

$$\begin{aligned} \min_{\substack{x_1, \dots, x_N \\ u_1, \dots, u_N}} \quad & \sum_{i=1}^N \left(\sum_{\tau=0}^{T-1} \ell_i^{SC}(x_i(\tau), u_i(\tau)) + \ell_i^{TC}(x_i(T)) \right) \\ \text{subj. to } \quad & x_i(\tau+1) = A_i x_i(\tau) + B_i u_i(\tau), \quad \forall i, \quad \tau \in [0, T-1] \\ & x_i(\tau) \in X_i \quad u_i(\tau-1) \in U_i, \quad \forall i, \quad \tau \in [1, T] \\ & \sum_{i=1}^N H_i x_i(\tau) \leq \bar{h} \quad \tau \in [1, T] \end{aligned}$$

12.4 Duality for constraint-coupled optimization

Let the Lagrangian function $\mathcal{L} : \mathbb{R}^{\sum_{i=1}^N n_i} \times \mathbb{R}^r \rightarrow \mathbb{R}$ be

$$\mathcal{L}(x, \mu) = \sum_{i=1}^N (\ell_i(z_i) + \mu^T g_i(z_i))$$

Then, the dual function can be computed as

$$\begin{aligned}
 q(\mu) &= \inf_{\substack{z_i \in Z_i \\ i \in \{1, \dots, N\}}} \sum_{i=1}^N (\ell_i(z_i) + \mu^T g_i(z_i)) \\
 &= \sum_{i=1}^N \inf_{z_i \in Z_i} (\ell_i(z_i) + \mu^T g_i(z_i)) \\
 &= \sum_{i=1}^N q_i(\mu)
 \end{aligned}$$

with $q_i(\mu) := \inf_{z_i \in Z_i} (\ell_i(z_i) + \mu^T g_i(z_i))$. Because each local cost function depends only on its own variable and the constraint sets are independent we can minimize each cost function separately.

The dual problem is thus

$$\max_{\mu \geq 0} \sum_{i=1}^N q_i(\mu)$$

Remark. The dual problem is a cost-coupled optimization problem.

12.5 Centralized dual subgradient

The subgradient algorithm non the dual problem reads

$$\begin{aligned}
 \mu^{k+1} &= \mathcal{P}_{\mu \geq 0} \left(\mu^k + \alpha^k \tilde{\nabla} q(\mu^k) \right) \\
 &= \mathcal{P}_{\mu \geq 0} \left(\mu^k + \alpha^k \sum_{i=1}^N \tilde{\nabla} q_i(\mu^k) \right)
 \end{aligned}$$

where α^k is the stepsize.

A subgradient of q_i at μ^k can be computed as

$$\begin{aligned}
 z_i^{k+1} &= \arg \min_{z_i \in Z_i} \ell_i(z_i) + (\mu^k)^T g_i(z_i) \\
 \tilde{\nabla} q_i(\mu^k) &= g_i(z_i^{k+1})
 \end{aligned}$$

Centralized dual subgradient algorithm

$$\begin{aligned}
 z_i^{k+1} &= \arg \min_{z_i \in Z_i} \ell_i(z_i) + (\mu^k)^T g_i(z_i) \\
 \mu^{k+1} &= \mathcal{P}_{\mu \geq 0} \left(\mu^k + \alpha^k \sum_{i=1}^N g_i(z_i^{k+1}) \right)
 \end{aligned}$$

Remark. The centralized part is the update of μ , which is to be computed by a "master" node, whereas the update of z_i can be done locally to each agent.

Remark. This procedure is known as dual decomposition, as the problem is decomposed and computation can be parallelized

Distributed dual subgradient algorithm

$$\begin{aligned}
 v_i^{k+1} &= \sum_{j=1}^N a_{ij} \mu_j^k \\
 z_i^{k+1} &= \arg \min_{z_i \in Z_i} \ell_i(z_i) + (v_i^{k+1})^T g_i(z_i) \\
 \mu_i^{k+1} &= \mathcal{P}_{z \geq 0} (\mu_i^k + \alpha^k g_i(z_i^{k+1}))
 \end{aligned}$$

This algorithm was obtained by applying the distributed gradient algorithm to the centralized algorithm, thus obtaining a completely distributed algorithm.

If the minimum in the second step can be proven to be unique at every iteration, then z_i^k converges to the optimal solution of the primal problem. Otherwise, it can be proven that the running average of z_i^k converges to the optimal solution.

12.5.1 Distributed dual subgradient convergence

Theorem 12.1 (Distributed Dual Subgradient)

Let Assumptions 1, 2, and 3d hold.